



# F A S T   S c h o o l   o f   C o m p u t i n g

## CS4064 – Advance Database Concepts

**Spring 2025**

**Instructor Name:** Muhammad Ishaq Raza  
**Email address:** ishaq.raza@nu.edu.pk  
**Office Location/Number:** NB-F-017  
**Office Hours:** Tu, We, Th 1:30 - 2:30 PM

**TA Name (BCS-6A):** Siftullah  
**Email:** l215263@lhr.nu.edu.pk  
**TA Name (BCS-6B):** Saleh Ahmed  
**Email:** l211826@lhr.nu.edu.pk

### Course Information

**Program:** BS      **Credit Hours:** 3      **Type:** Elective

**Pre-requisites:** CS2005 – Database Systems

### Course Learning Outcomes (CLOs)

This course is intended for students who wish to specialize in database management systems or wish to practice the advanced techniques involved in optimization of data storage, database design and queries. The course primarily addresses design and implementation of a database management system. It covers advanced topics like physical storage and access methods, query optimization, transaction processing, concurrency control, and recovery techniques.

CLO#	CLO Description
CLO 1	<b>Understanding</b> advance data models, technologies, and approaches for building database systems.
CLO 2	<b>Apply</b> the models and approaches in order to become enabled to select and apply appropriate methods for a particular case.
CLO 3	To <b>develop</b> a database solution for given scenario/challenging problem in the domain of database systems.

### Textbook

- Ramez Elmasri, Fundamentals of Database Systems (7<sup>th</sup> Edition)

### Reference Books

- Raghu Ramakrishnan, *Database Management Systems* (3<sup>rd</sup> Edition)
- Jeffrey D. Ullman, Jennifer Widom, Hector Garcia-Molina, *Database System Implementation*
- George Coulouris, *Distributed Systems; Concepts and Design* (3<sup>rd</sup> Edition)

### Grading Scheme (Absolute)

Midterms (30%), Quizzes (10%), Assignments/Class Participation (10%), Final (50%)

### Grading

- Minimum eligibility to pass this course is to get 50% marks.
- No makeup for missed quizzes or assignments and quizzes may be un-announced.
- Academic integrity is expected of all the students. Plagiarism or cheating in any assessment may result in an **F** grade in the course, and possibly more severe penalties.

## Tentative Weekly Schedule

Week	Topics to be covered	Topics Detail	Readings (Textbook)	No of Lec.	Asst.
1-2	Transaction Processing Concepts and Theory	<ul style="list-style-type: none"> <li>Issues in Transaction Processing</li> <li>Why Concurrency Control is Needed</li> <li>Why Recovery is Needed</li> <li>Transaction States and Operations, System Log, Commit Point of a Transaction</li> <li>ACID Properties of Transactions</li> <li>Characterizing Schedules based on Recoverability</li> <li>Characterizing Schedules based on Serializability</li> <li>Transactions Isolation Levels and Possible Violations</li> </ul>	Ch 20	3	
2-3	Concurrency Control Techniques	<ul style="list-style-type: none"> <li>Two-phase locking techniques for concurrency control</li> <li>Concurrency control based on timestamp ordering</li> <li>Multiversion concurrency control techniques</li> <li>Validation (optimistic) techniques and snapshot isolation concurrency control</li> </ul>	Ch 21	3	A1
4-5	Database Recovery Techniques	<ul style="list-style-type: none"> <li>Recovery concepts</li> <li>NO-UNDO/REDO recovery based on deferred update</li> <li>Recovery techniques based on immediate update</li> <li>Shadow paging</li> <li>The ARIES recovery algorithm</li> <li>Recovery in multidatabase systems</li> <li>Database backup and recovery from catastrophic failures</li> </ul>	Ch 22	3	
5-7	Data Storage, Basic File Structures and Hashing	<ul style="list-style-type: none"> <li>Disk storage devices</li> <li>Buffering of blocks</li> <li>Placing file records on disk</li> <li>Operations on files</li> <li>Files of unordered records (heap files)</li> <li>Files of ordered records (sorted files)</li> <li>Hashing techniques: internal hashing, static hashed files (external hashing)</li> <li>Hashing techniques that allow dynamic file expansion: dynamic, extendible, and linear hashing techniques</li> </ul>	Ch 16	4	A2
7-10	Indexing Structures for Files and Physical Database Design	<ul style="list-style-type: none"> <li>Type of single-level ordered indexes: primary indexes, clustering indexes, secondary indexes</li> <li>Multilevel indexes: two level indexes, multilevel indexes</li> <li>Dynamic multilevel indexes using B-Trees, B<sup>+</sup>-Trees and B<sup>+</sup>*-Trees</li> <li>Indexes on multiple keys: ordered indexes on multiple attributes, partitioned hashing, grid files</li> <li>Other types of indexes: hash indexes, bitmap indexes, function-based indexing</li> <li>Physical database design in relational databases</li> <li>An overview of database tuning in relational systems</li> </ul>	Ch 17	6	A3
10-11	Strategies for Query Processing	<ul style="list-style-type: none"> <li>Translating SQL queries into relational algebra and other operators</li> <li>Algorithms for external sorting</li> <li>Algorithms for SELECT operation</li> <li>Algorithms for JOIN operation</li> <li>Algorithms for PROJECT and SET operations</li> <li>Implementing aggregate operations and different type of JOINS</li> <li>Combining operations using pipelining</li> <li>Parallel algorithms for query processing: parallel database architecture approaches (share-memory, shared-disk, and shared-nothing architectures), operator level parallelism, intra-query parallelism, inter-query parallelism</li> </ul>	Ch 18	3	A4
12	Query Optimization	<ul style="list-style-type: none"> <li>Query trees and heuristics for query optimization</li> <li>Choices of query execution plans, nested subquery optimization, materialized views</li> <li>Use of selectivities in cost-based optimization</li> <li>Cost components for query optimization</li> <li>Catalog information used in cost functions</li> <li>Histograms</li> <li>Multi-relation queries and JOIN ordering choices (left-deep join tree, right-deep join tree, and bushy join tree)</li> <li>Adaptive and semantic query optimization</li> </ul>	Ch 19	2	
13	Distributed Database Concepts	<ul style="list-style-type: none"> <li>Distributed computing systems and Big data technologies</li> <li>Distributed database concepts; Data fragmentation, replication, and allocation techniques for distributed database design</li> <li>Overview of transaction management in distributed databases</li> <li>Types of distributed database systems</li> <li>Distributed database architectures.</li> </ul>	Ch 23	2	
14	NOSQL Databases and Big Data Storage Systems	<ul style="list-style-type: none"> <li>Introduction to NOSQL systems</li> <li>The CAP theorem</li> <li>Document-based NOSQL systems and MongoDB; NOSQL Key-Value stores; column-based NOSQL systems; NOSQL graph databases and Neo4j</li> </ul>	Ch 24	2	